

→ Attention: multi-head attention

- Q, K, V are all the input X .
- so we take $X \rightarrow (\text{batch-len}, \text{seq-len}, \text{d-model})$ and project them into a different dimension.
- a linear model will be used to project them.
 - model $(\text{d-model}, \text{d-model})$
 - out-q: model (Q)
 - then transform out-q $\rightarrow q\text{-proj}$
 - $(\text{batch-len}, \text{seq-len}, \text{d-model}) \rightarrow (\text{batch-len}, \text{seq-len}, \text{num-heads}, \text{d-model})$
- Transform from $q\text{-proj}(\text{batch-len}, \text{seq-len}, \text{num-heads}, \text{d-model})$ to $q\text{-proj}(\text{batch-len}, \text{num-heads}, \text{seq-len}, \text{d-model})$
- then perform scaled dot product attention

$$\frac{Q \cdot K^T}{\sqrt{d_k}} \rightarrow \text{then apply mask}$$

(essentially fill the indices in mask with a '-∞')

↓ apply softmax

↓ @V → result.

→ the result from scaled dot product attention is of shape $(\text{batch-size}, \text{num-heads}, \text{seq-len}, \text{d-model})$

↓ transpose

batch-size, seq-len, num-heads, d-model

↓ apply out model put

↓ reshape output to $(\text{batch-len}, -1, \text{d-model})$

→ Positional encodings:

$$PE = \begin{cases} \sin(\text{pos}/1000^{2i/d\text{-model}}) \\ \cos(\text{pos}/1000^{2i/d\text{-model}}) \end{cases}$$

create a pe matrix where terms must be

div-term will produce

$$\left[\frac{1}{1000^{2 \cdot 0/d}}, \frac{1}{1000^{2 \cdot 1/d}}, \dots, \frac{1}{1000^{2 \cdot (d/2)/d}} \right]$$

positions will be

$$\begin{bmatrix} 0 \\ \vdots \\ L-1 \end{bmatrix}$$

$d \rightarrow d\text{-model}$
 $L \rightarrow \text{seq-len}$

$$PE = \text{batch-size} \times \text{pos}(\text{positions} \times \text{div})$$

$$PE = \begin{bmatrix} \sin(0) & \cos(0) & \sin(2) & \dots & \dots \\ \sin(1) & \cos(1) & \dots & \dots & \dots \\ \sin(2) & \dots & \dots & \dots & \dots \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ \sin(L-1) & \dots & \dots & \dots & \dots \end{bmatrix}$$

div-terms are made using the formula

$$\exp(2i \cdot \frac{-\log(1000)}{d}) \approx \frac{1}{1000^{2i/d\text{-model}}}$$

→ Layer norm:

$$\frac{x - \text{mean}}{\sqrt{\text{var} + \text{eps}}} * \gamma + \beta$$

trainable parameters

Layer norm will only applied on the -1 dimension, which is each layer.

→ Batch Normalization:


$$\frac{x - \text{mean}}{\sqrt{\text{var} + \text{eps}}} * \gamma + \beta$$

applied on each batch on 0 dimension.

→ Transformer:


attention(X, X, X, mask) → dropout

dropout is done before 2 layer norms.



→ Language Model:

embeddings → positional encodings → create a mask (for attention)



* before passing inputs to the language model, we need to convert the inputs into tokens.

→ Tokenization:

normal tokenization:

- convert the text by breaking down words to characters including all the symbols. (data sample level)
- now store all the unique characters.
- sort these unique chars to form a vocabulary. and len of the vocabulary will be the vocab size.
- map the indices of characters in the form below
 - $\{ch: i\}$ → character in the vocabulary.
- now make index using the $\{ch: i\}$ mapping.
- store id's of the characters present in each sample by mapping them to the stoi.
- make pairs :- suppose we have ids for chars like c, h, a, r
 - 0, 1, 2, 3
 - it creates pairs like $\{0, 1, 2, 3\}$
 - $\{(0, 1), (2, 3)\}$ $\{(0, 1, 2), (3)\}$
- now initialize a seq-len
 - use a sliding window approach to convert the data into chunks
 - all-ids ← all input ids from every sample.
 - so each chunk will have
 - $x = \text{all-ids}[i : i + \text{seq-len}]$
 - $y = \text{all-ids}[i+1 : i + \text{seq-len} + 1]$
 - input-ids will be x and labels will be y.
- this is all for training set.
- for testing set, we don't want to create a stoi, instead use the stoi from training.
- vocab-size will automatically be the len of stoi.

→ BPE tokenizer:

Byte pair encoding tokenizer

- maintain a dict
 - first given the text
 - convert the characters into bytes (hex) and then to list
 - list(`Text.encode('utf-8')`)
 - each char takes 1 bytes and some complex char like emoji's take up 4 bytes.
- we will have to specify a vocab size. so BPE only uses 256 ids. hyper-parameter.
- if hyper-parameter is 300, then 44 merges must happen.
- create pairs and generate a counter for the pairs.
- get the most freq pair and now use a next-id store the next-id in token-to-char with value of concatenation of the pairs.
- create a variable which tracks the merges.
- increment next-id + 1.

→ encoding using BPE:

essentially comes down to merging pairs stored in merges.

merging pairs →

create a new token list where you check

single iteration { for the occurrence of the pair, if there is one instead of sending in pair (0) and pair (1) into list, we send in the next id (both a and b)

for each merge pair in merges, run this iteration.

→ decode:

for token into tokens:

append the utf decoded value of the token using utf.decode and the bytes stored in token-to-char.

if token is not recognizable, add empty char to the text.

→ Training losses:

- scale → torch.nn.CrossEntropyLoss
- apply it on loss, optimizer.
- torch.nn.CrossEntropyLoss
- use it for calculating logits

→ Metrics:

Perplexity = $\text{math.exp}(\text{avg-loss})$

↓ how confident model is while making predictions.

→ generate text:

* first generate logits, and scale by temperature

logits[:, -1, :]/temperature.

* instead of using argmax to obtain highest probs. use multinomial sampling.

* next-token-id = torch.multinomial(probs, num_samples=1)

* concatenate the next-token-id with x.

run this iteration for max-new-token times.

at last use tokenizer.decode to convert into text.

→ difference between multinomial Sampling vs argmax

be → 0.7

banana → 0.1

to → 0.2

argmax picks be since high prob score

but multinomial sampling picks randomly b/w three, like a dice rolling with three probable outcomes.